

# Operating Systems (AI 3002) – 120+ Viva Q&A

## Contents

<b>1</b>	<b>Section I: Introduction to Operating System</b>	<b>2</b>
1.1	Unit 1: Introduction to OS (4 Hours) . . . . .	2
1.2	Unit 2: Process Management (6 Hours) . . . . .	3
<b>2</b>	<b>Section II: CPU Scheduling, Deadlock, Memory, and I/O</b>	<b>5</b>
2.1	Unit 3: Uniprocessor Scheduling (4 Hours) . . . . .	5
2.2	Unit 4: Deadlock (4 Hours) . . . . .	6
2.3	Unit 5: Memory Management (5 Hours) . . . . .	7
2.4	Unit 6: I/O and File Management (5 Hours) . . . . .	8
<b>3</b>	<b>Practicals, AI Assignments, and Projects</b>	<b>10</b>

# SECTION I: INTRODUCTION TO OPERATING SYSTEM

## Unit 1: Introduction to OS (4 Hours)

**Q1.** What is an Operating System?

**Ans:** An OS is system software that manages hardware resources and provides a platform for application programs to run. It acts as an intermediary between user and hardware.

**Q2.** List the primary goals of an OS.

**Ans:** Convenience (user-friendly interface), efficiency (optimal resource utilization), and ability to evolve (support new functions without interference).

**Q3.** Name five services provided by an OS.

**Ans:** Program execution, I/O operations, file system manipulation, communication, error detection.

**Q4.** What are system calls? Give five examples.

**Ans:** System calls are kernel-level interfaces for user programs. Examples: `fork()`, `exec()`, `wait()`, `open()`, `read()`.

**Q5.** Differentiate between system calls and library functions.

**Ans:** System calls involve kernel mode switch and OS services; library functions run in user space and may or may not use system calls.

**Q6.** What is a batch OS? Give one drawback.

**Ans:** Jobs are grouped and executed without user interaction. Drawback: No interactivity; low CPU utilization due to I/O waits.

**Q7.** Explain multiprogramming.

**Ans:** Multiple programs reside in memory; CPU switches among them to keep itself busy while one waits for I/O.

**Q8.** What is time-sharing? How is it different from multiprogramming?

**Ans:** Time-sharing allows multiple users to interact simultaneously using CPU time slices. Multiprogramming focuses on CPU utilization; time-sharing adds interactivity.

**Q9.** Define real-time OS. Give two types.

**Ans:** An OS that guarantees task completion within strict time limits. Types: Hard RTOS (missed deadline = system failure) and Soft RTOS (occasional misses acceptable).

**Q10.** What is a distributed OS?

**Ans:** An OS that manages a collection of independent computers appearing as a single system to users. Resources are shared across networked machines.

**Q11.** What is a parallel OS?

**Ans:** An OS designed to manage multiprocessor systems where multiple CPUs work simultaneously on a single task.

**Q12.** Explain the evolution from batch to modern OS.

**Ans:** Batch (no interaction) → Multiprogramming (better CPU use) → Time-sharing (interactivity) → Personal computing → Networked/distributed → Mobile/cloud OS.

**Q13.** What is BASH?

**Ans:** Bourne Again SHell—a command-line interpreter for Unix/Linux that supports scripting and automation.

**Q14.** Give five basic Linux commands.

**Ans:** ls, cd, pwd, mkdir, rm.

**Q15.** What is the role of a shell in OS?

**Ans:** A shell is a command interpreter that provides a user interface to access OS services via commands or scripts.

## Unit 2: Process Management (6 Hours)

**Q16.** Define a process.

**Ans:** A program in execution, including its code, data, stack, heap, and PCB (Process Control Block).

**Q17.** What is a Process Control Block (PCB)? What does it contain?

**Ans:** A data structure storing process state: PID, program counter, CPU registers, memory limits, open files, scheduling info, etc.

**Q18.** Draw and explain the 5-state process model.

**Ans:** New → Ready → Running → (Blocked/Waiting) → Terminated. Blocked occurs on I/O wait; Ready means ready to run but not scheduled.

**Q19.** What is context switching? Why is it costly?

**Ans:** Saving the state of one process and loading another. Costly due to CPU overhead (cache flush, TLB invalidation, register save/restore).

**Q20.** What is a thread? Why use threads?

**Ans:** A lightweight unit of execution within a process. Advantages: Faster creation, shared memory, improved responsiveness, better utilization of multiprocessors.

**Q21.** Differentiate user-level and kernel-level threads.

**Ans:** User threads: managed by user library; kernel unaware → faster but blocking one blocks all. Kernel threads: managed by OS → true concurrency but slower.

**Q22.** What is concurrency? What problems arise due to it?

**Ans:** Multiple computations executing during overlapping time periods. Problems: race conditions, deadlock, starvation, inconsistent data.

**Q23.** What is a race condition?

**Ans:** When outcome depends on the sequence/timing of uncontrollable events (e.g., two threads modifying shared variable without sync).

**Q24.** Define critical section.

**Ans:** A code segment accessing shared resources that must not be concurrently executed by more than one process/thread.

**Q25.** State the three requirements for solving critical section problem.

**Ans:** Mutual exclusion, progress, bounded waiting.

**Q26.** What is a mutex? How is it implemented?

**Ans:** A mutual exclusion lock allowing only one thread to enter critical section. Implemented via atomic instructions like Test-and-Set or Compare-and-Swap.

**Q27.** What is a semaphore? Types?

**Ans:** A synchronization variable. Binary (0/1) for mutual exclusion; counting ( $\geq 1$ ) for resource counting.

**Q28.** Explain the Producer-Consumer problem.

**Ans:** Producer adds items to buffer; consumer removes. Use empty/full semaphores and mutex to avoid overflow/underflow.

**Q29.** Describe the Readers-Writers problem.

**Ans:** Multiple readers can access data simultaneously, but writers need exclusive access. Solutions prioritize readers or writers.

**Q30.** Explain the Dining Philosophers problem.

**Ans:** Five philosophers alternate between thinking and eating; each needs two adjacent forks. Poor coordination leads to deadlock or starvation.

**Q31.** How to solve Dining Philosophers using semaphores?

**Ans:** Use an array of 5 mutexes for forks and a global semaphore limiting concurrent eaters to 4, breaking circular wait.

**Q32.** What is a monitor?

**Ans:** A high-level synchronization construct that encapsulates shared data and procedures, ensuring only one process executes inside at a time.

**Q33.** Why are orphan and zombie processes created?

**Ans:** Zombie: child terminates but parent hasn't read its exit status. Orphan: parent dies before child; child adopted by init.

**Q34.** How does `fork()` work?

**Ans:** Creates a child process as a copy of the parent. Returns 0 to child, child PID to parent.

**Q35.** What is `exec()`?

**Ans:** Replaces current process image with a new program (e.g., `execvp("ls", args)`).

**Q36.** What is `wait()`?

**Ans:** Parent suspends until child terminates, then collects exit status to prevent zombie.

**Q37.** Explain Pthreads.

**Ans:** POSIX threads API for multithreading in C. Includes functions like `pthread_create()`, `pthread_join()`, `pthread_mutex_lock()`.

**Q38.** What is thread safety?

**Ans:** Code that functions correctly when accessed by multiple threads simultaneously (uses proper synchronization).

**Q39.** What is deadlock in multithreading?

**Ans:** Two or more threads wait indefinitely for resources held by each other (e.g., Thread A locks X, waits for Y; Thread B locks Y, waits for X).

**Q40.** How to avoid race conditions in threads?

**Ans:** Use mutexes, semaphores, atomic operations, or thread-safe data structures.

## SECTION II: CPU SCHEDULING, DEADLOCK, MEMORY, AND I/O

### Unit 3: Uniprocessor Scheduling (4 Hours)

**Q41.** What is CPU scheduling?

**Ans:** The process of selecting which ready process gets the CPU next.

**Q42.** List scheduling criteria.

**Ans:** CPU utilization, throughput, turnaround time, waiting time, response time.

**Q43.** What is turnaround time?

**Ans:** Total time from submission to completion: Completion - Arrival.

**Q44.** Define waiting time.

**Ans:** Total time spent in ready queue.

**Q45.** Explain FCFS scheduling. Is it preemptive?

**Ans:** First-Come-First-Served; non-preemptive. Suffers from convoy effect.

**Q46.** What is SJF? Why is it optimal?

**Ans:** Shortest Job First minimizes average waiting time if burst times are known.

**Q47.** Differentiate preemptive and non-preemptive SJF.

**Ans:** Non-preemptive: once started, runs to completion. Preemptive (SRTF): if new shorter job arrives, current is preempted.

**Q48.** How does Round Robin work?

**Ans:** Each process gets a fixed time quantum. After quantum, it's preempted and moved to end of ready queue.

**Q49.** What happens if time quantum is too small or too large in RR?

**Ans:** Too small → high context switch overhead. Too large → behaves like FCFS.

**Q50.** Explain Priority Scheduling. What is starvation?

**Ans:** Highest priority process runs first. Starvation: low-priority processes may never run.

**Q51.** How to prevent starvation in priority scheduling?

**Ans:** Aging: gradually increase priority of waiting processes.

**Q52.** Which algorithm is used in modern OS?

**Ans:** Multilevel Feedback Queue (not in syllabus but conceptually built from these), combining RR and priority with aging.

**Q53.** Can SJF be implemented in real life? Why/why not?

**Ans:** Rarely—burst times are unknown. Approximated using exponential averaging.

**Q54.** What is response time? Why is it important for interactive systems?

**Ans:** Time from request to first response. Critical for user experience in interactive apps.

**Q55.** Compare FCFS and RR in terms of fairness.

**Ans:** FCFS unfair to short jobs; RR fair as all get equal CPU time.

#### **Unit 4: Deadlock (4 Hours)**

**Q56.** Define deadlock.

**Ans:** A situation where a set of processes are blocked forever, each waiting for a resource held by another.

**Q57.** State four necessary conditions for deadlock.

**Ans:** Mutual exclusion, hold and wait, no preemption, circular wait.

**Q58.** How does mutual exclusion contribute to deadlock?

**Ans:** Only one process can use a resource at a time; others must wait.

**Q59.** Explain “hold and wait.”

**Ans:** A process holds at least one resource while waiting for others.

**Q60.** How to prevent hold-and-wait?

**Ans:** Require processes to request all resources before execution begins.

**Q61.** What is deadlock prevention?

**Ans:** Ensuring at least one of the four necessary conditions never holds.

**Q62.** What is deadlock avoidance?

**Ans:** Dynamically examine resource allocation state to ensure it never enters an unsafe state (e.g., Banker’s algorithm).

**Q63.** Explain Banker’s algorithm.

**Ans:** Before allocating, check if system remains in safe state (all processes can finish in some order).

**Q64.** What is a safe state? Unsafe state?

**Ans:** Safe: there exists a safe sequence. Unsafe: may lead to deadlock.

**Q65.** What is deadlock detection?

**Ans:** Allow deadlock to occur; use algorithms (e.g., resource allocation graph) to detect and recover.

**Q66.** How does resource allocation graph detect deadlock?

**Ans:** Cycle in graph with single-instance resources → deadlock.

**Q67.** What is deadlock recovery? Methods?

**Ans:** Abort processes or preempt resources. Methods: kill all deadlocked, kill one at a time, rollback to checkpoint.

**Q68.** Why is Banker’s algorithm rarely used in real systems?

**Ans:** Requires fixed number of processes/resources; assumes known max demand—unrealistic in dynamic environments.

**Q69.** Can deadlock occur with reusable resources only?

**Ans:** Yes—e.g., memory, files, printers.

**Q70.** Can two processes deadlock with one resource?

**Ans:** No—circular wait requires at least two resources.

**Q71.** How does Dining Philosophers illustrate deadlock?

**Ans:** All pick up left fork simultaneously → each waits for right fork → circular wait → deadlock.

## Unit 5: Memory Management (5 Hours)

**Q72.** What are memory management requirements?

**Ans:** Relocation, protection, sharing, logical organization, physical organization.

**Q73.** What is static vs dynamic loading?

**Ans:** Static: load entire program at start. Dynamic: load parts on demand → saves memory.

**Q74.** Explain fixed and variable partitioning.

**Ans:** Fixed: memory divided into fixed partitions (internal fragmentation). Variable: partitions created per process (external fragmentation).

**Q75.** What is internal fragmentation?

**Ans:** Wasted space within allocated memory block (e.g., in paging).

**Q76.** What is external fragmentation?

**Ans:** Free memory exists but not in contiguous blocks (e.g., in segmentation). Solved by compaction or paging.

**Q77.** What is paging?

**Ans:** Memory divided into fixed-size frames; program into pages. Enables non-contiguous allocation.

**Q78.** How is logical address converted to physical in paging?

**Ans:** Logical = page# + offset; use page table to get frame#; physical = frame# + offset.

**Q79.** What is a page table?

**Ans:** Maps page numbers to frame numbers in physical memory.

**Q80.** Explain segmentation.

**Ans:** Memory divided into variable-size segments (code, data, stack). Logical address = segment# + offset.

**Q81.** Compare paging and segmentation.

**Ans:** Paging: fixed size, no external frag, internal frag. Segmentation: logical units, external frag, no internal frag.

**Q82.** What is virtual memory?

**Ans:** Technique to execute processes larger than physical memory using disk as extension.

**Q83.** How is virtual memory implemented?

**Ans:** Via demand paging—pages loaded only when referenced.

**Q84.** What is a page fault?

**Ans:** Occurs when referenced page is not in RAM; OS loads it from disk.

**Q85.** Explain FIFO page replacement.

**Ans:** Replace oldest loaded page. Suffers from Belady's anomaly.

**Q86.** What is Belady's anomaly?

**Ans:** Increasing page frames may increase page faults (in FIFO).

**Q87.** Explain LRU.

**Ans:** Replace least recently used page. Approximates optimal but costly to implement.

**Q88.** What is Optimal page replacement?

**Ans:** Replace page not used for longest future time. Theoretical minimum faults; not implementable (needs future knowledge).

**Q89.** What is TLB? Why used?

**Ans:** Translation Lookaside Buffer—cache for recent page table entries to speed up address translation.

**Q90.** What is thrashing? How to prevent?

**Ans:** High page fault rate due to insufficient frames. Prevent via working-set model or page fault frequency control.

**Q91.** Explain placement strategies: First Fit, Best Fit, Worst Fit.

**Ans:** First Fit: allocate first hole size. Best Fit: smallest sufficient hole. Worst Fit: largest hole. Best Fit → small unusable holes; Worst Fit → less fragmentation.

**Q92.** Why is paging preferred over segmentation in modern OS?

**Ans:** Simpler hardware, no external fragmentation, easier swapping.

## Unit 6: I/O and File Management (5 Hours)

**Q93.** Why is I/O management complex?

**Ans:** Devices vary in speed, data rate, protocol; require buffering, error handling, and scheduling.

**Q94.** What is I/O buffering? Why used?

**Ans:** Temporary storage to absorb speed mismatches between CPU and devices.

**Q95.** Types of I/O buffering?

**Ans:** Single, double, circular buffering.

**Q96.** What are disk scheduling algorithms?

**Ans:** Methods to order disk read/write requests to minimize seek time.

**Q97.** Explain FCFS disk scheduling.

**Ans:** Serve requests in arrival order. Simple but inefficient.

**Q98.** What is SSTF? Drawback?

**Ans:** Shortest Seek Time First—serve nearest request. Drawback: starvation of distant requests.

**Q99.** Explain SCAN algorithm.

**Ans:** Disk arm moves in one direction, serving all requests; reverses at end.

**Q100.** What is C-SCAN?

**Ans:** Circular SCAN—arm moves to end, then jumps to start without servicing—uniform wait time.

**Q101.** Compare SCAN and C-SCAN.

**Ans:** SCAN has lower average seek time; C-SCAN gives more uniform performance.

**Q102.** What is a file?

**Ans:** Logical collection of related data with a name, stored on secondary storage.

**Q103.** What is a file system?

**Ans:** Mechanism to store, organize, and retrieve files (e.g., ext4, NTFS).

**Q104.** What is file organization? Types?

**Ans:** How records are stored: sequential, indexed, hashed.

**Q105.** What is sequential access? Direct access?

**Ans:** Sequential: read in order. Direct: access any record via offset/key.

**Q106.** What is a directory?

**Ans:** File containing metadata (names, attributes, locations) of other files.

**Q107.** Explain single-level, two-level, tree-structured directories.

**Ans:** Single: all files in one dir. Two-level: user + file. Tree: hierarchical (modern OS).

**Q108.** What is file sharing? Problems?

**Ans:** Multiple users access same file. Problems: consistency, access control, concurrency.

**Q109.** What is record blocking?

**Ans:** Storing multiple records in one disk block to reduce I/O operations.

**Q110.** What is secondary storage management?

**Ans:** OS functions: disk formatting, free-space management (bit vector, linked list), allocation (contiguous, linked, indexed).

**Q111.** Compare contiguous, linked, and indexed allocation.

**Ans:** Contiguous: fast access, external frag. Linked: no frag, slow access. Indexed: supports direct access, overhead of index block.

**Q112.** What is FAT?

**Ans:** File Allocation Table—linked allocation with table in memory for fast access (used in MS-DOS).

**Q113.** What is inode?

**Ans:** Data structure in Unix storing file metadata and block pointers.

**Q114.** What is spooling?

**Ans:** Simultaneous Peripheral Operations On-Line—buffering jobs (e.g., print queue) for devices slower than CPU.

**Q115.** How does OS handle I/O errors?

**Ans:** Retry, report to user, log error, or terminate process.

**Q116.** What is device independence?

**Ans:** Programs access devices via logical names (e.g., “printer”)—OS maps to physical device.

## PRACTICALS, AI ASSIGNMENTS, AND PROJECTS

**Q117.** How do you create a zombie process in C?

**Ans:** Parent calls `fork()`, child exits, parent sleeps without calling `wait()`.

**Q118.** How to avoid zombie processes?

**Ans:** Parent must call `wait()` or `waitpid()`. Alternatively, use double fork.

**Q119.** Write a shell script to list files with .txt extension.

**Ans:** `ls *.txt`

**Q120.** How to implement Banker’s algorithm in C?

**Ans:** Input: available, max, allocation → compute need → check safe sequence using work and finish arrays.

**Q121.** How to simulate FIFO page replacement?

**Ans:** Use queue—on page fault, if full, dequeue oldest and enqueue new.

**Q122.** What ML model can predict best CPU scheduler?

**Ans:** Decision tree or random forest trained on features like burst time, arrival pattern, I/O ratio.

**Q123.** How can LSTM predict page faults?

**Ans:** Train on page access sequences; LSTM learns temporal patterns to forecast next access.

**Q124.** What data is needed for deadlock prediction model?

**Ans:** Resource allocation state, process requests, hold patterns over time.

**Q125.** What is the role of system logs in anomaly detection?

**Ans:** Logs reveal unusual sequences (e.g., repeated failed logins)—autoencoders detect deviations.

**Q126.** What are key components of your multiprogramming OS project?

**Ans:** CPU simulator, PCB, scheduler, memory manager (paging), interrupt handler, IPC.

**Q127.** How is virtual memory implemented in your project?

**Ans:** Demand paging with page table, page fault handler, and disk swap space simulation.

**Q128.** What is inter-process communication (IPC)? Methods?

**Ans:** Mechanisms for processes to exchange data: pipes, message queues, shared memory, sockets.

Created by Anshul